

# Impala: A Modern, Open-Source SQL Engine for Hadoop

Mark Grover

Software Engineer, Cloudera

January 7th, 2014

Twitter: [mark\\_grover](#)

[github.com/markgrover/impala-thug/](https://github.com/markgrover/impala-thug/)



# Agenda

---

- What is Impala and what is Hadoop?
- Execution frameworks on Hadoop – MR, Hive, etc.
- Goals and user view of Impala
- Architecture of Impala
- Comparing Impala to other systems
- Impala Roadmap

# What is Impala?

---

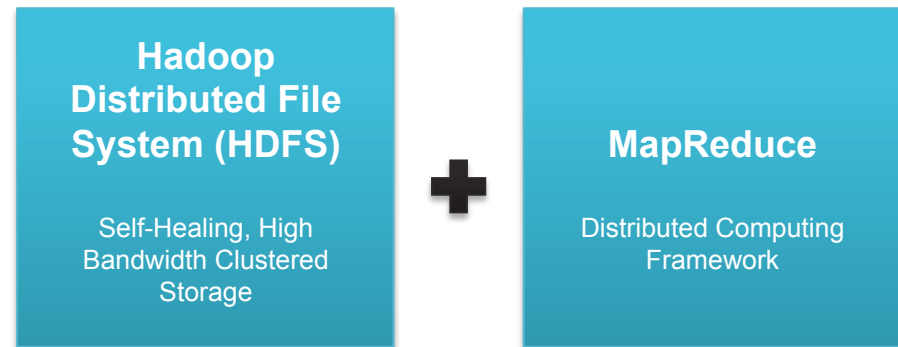
- General-purpose SQL engine
- Real-time queries in Apache Hadoop

# What is Apache Hadoop?

**Apache Hadoop** is an open source platform for data storage and processing that is...

- ✓ Distributed
- ✓ Fault tolerant
- ✓ Scalable

## CORE HADOOP SYSTEM COMPONENTS



### Has the Flexibility to Store and Mine Any Type of Data

- Ask questions across structured and unstructured data that were previously impossible to ask or solve
- Not bound by a single schema

### Excels at Processing Complex Data

- Scale-out architecture divides workloads across multiple nodes
- Flexible file system eliminates ETL bottlenecks

### Scales Economically

- Can be deployed on commodity hardware
- Open source platform guards against vendor lock

# So, what's wrong with MapReduce?

---

- Batch oriented
- High latency
- Not all paradigms fit very well
- Only for developers

# What are Hive and Pig?

---

- MR is hard and only for developers
- Higher level platforms for converting declarative syntax to MapReduce
  - SQL – Hive
  - workflow language – Pig
- Build on top of MapReduce

# What is Impala?

---

- General-purpose SQL engine
- Real-time queries in Apache Hadoop
- Beta version released since October 2012
- General availability (v1.0) release out since April 2013
- Open source under Apache license
- Latest release (v1.2.3) released on December 23rd

# Impala Overview: Goals

---

- General-purpose SQL query engine:
  - Works for both for analytical and transactional/single-row workloads
  - Supports queries that take from milliseconds to hours
- Runs directly within Hadoop:
  - reads widely used Hadoop file formats
  - talks to widely used Hadoop storage managers
  - runs on same nodes that run Hadoop processes
- High performance:
  - C++ instead of Java
  - runtime code generation
  - completely new execution engine – No MapReduce



# User View of Impala: Overview

---

- Runs as a distributed service in cluster: one Impala daemon on each node with data
- Highly available: no single point of failure
- User submits query via ODBC/JDBC, Impala CLI or Hue to any of the daemons
- Query is distributed to all nodes with relevant data
- Impala uses Hive's metadata interface, connects to Hive metastore

# User View of Impala: Overview

---

- There is no 'Impala format'!
- There is no 'Impala format'!!
- Supported file formats:
  - uncompressed/lzo-compressed text files
  - sequence files and RCFile with snappy/gzip compression
  - Avro data files
  - Parquet columnar format (more on that later)

# User View of Impala: SQL

---

- SQL support:
  - essentially SQL-92, minus correlated subqueries
  - INSERT INTO ... SELECT ...
  - only equi-joins; no non-equi joins, no cross products
  - Order By requires Limit
  - (Limited) DDL support
  - SQL-style authorization via Apache Sentry (incubating)
  - UDFs and UDAFs are supported

# User View of Impala: SQL

---

- Functional limitations:
  - No file formats, SerDes
  - no beyond SQL (buckets, samples, transforms, arrays, structs, maps, xpath, json)
  - Broadcast joins and partitioned hash joins supported
    - Smaller table has to fit in aggregate memory of all executing nodes

# User View of Impala: HBase

---

- Functionality highlights:
  - Support for SELECT, INSERT INTO ... SELECT ..., and INSERT INTO ... VALUES(...)
  - Predicates on rowkey columns are mapped into start/stop rows
  - Predicates on other columns are mapped into SingleColumnValueFilters
- But: mapping of HBase tables metastore table patterned after Hive
  - All data stored as scalars and in ascii
  - The rowkey needs to be mapped into a single string column

# User View of Impala: HBase

---

- Roadmap
  - Full support for UPDATE and DELETE
  - Storage of structured data to minimize storage and access overhead
  - Composite row key encoding, mapped into an arbitrary number of table columns

# Impala Architecture

---

- Three binaries: impalad, statestored, catalogd
- Impala daemon (impalad) – N instances
  - handles client requests and all internal requests related to query execution
- State store daemon (statestored) – 1 instance
  - Provides name service and metadata distribution
- Catalog daemon (catalogd) – 1 instance
  - Relays metadata changes to all impalad's

# Impala Architecture

---

- Query execution phases
  - request arrives via odbc/jdbc
  - planner turns request into collections of plan fragments
  - coordinator initiates execution on remote impalad's




# Impala Architecture

---

- During execution
  - intermediate results are streamed between executors
  - query results are streamed back to client
  - subject to limitations imposed to blocking operators (top-n, aggregation)

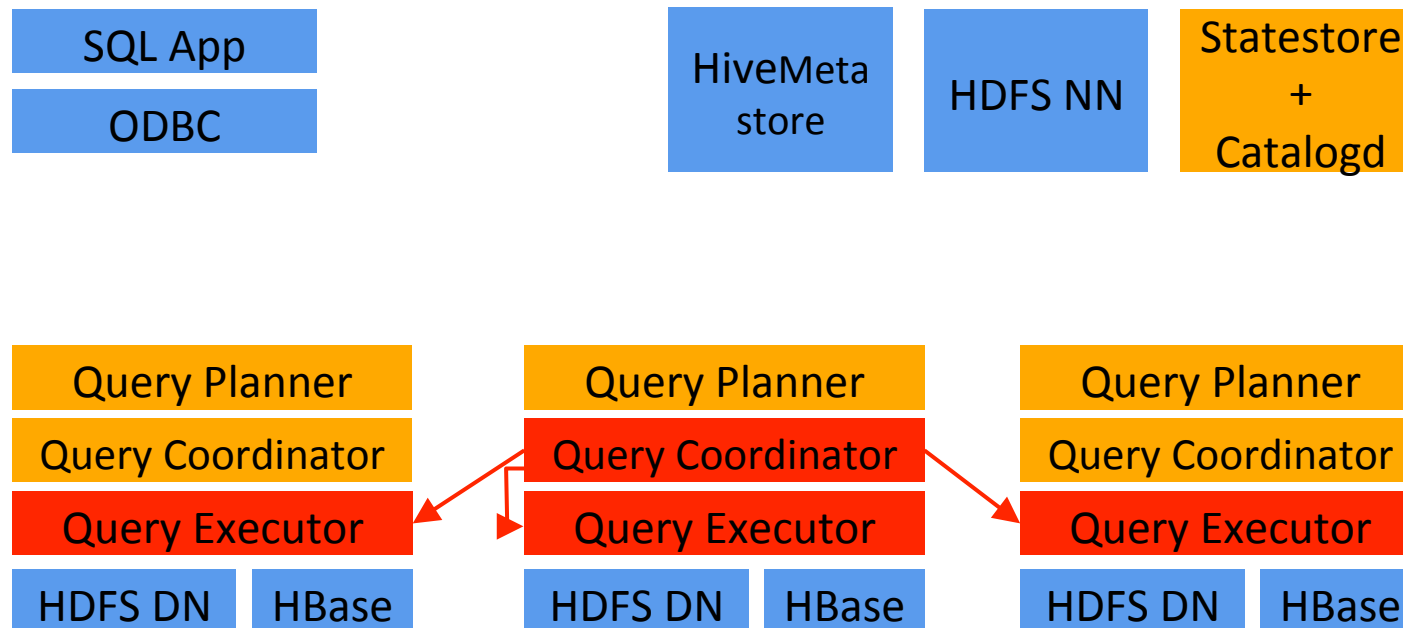
# Impala Architecture: Query Execution

Request arrives via odbc/jdbc



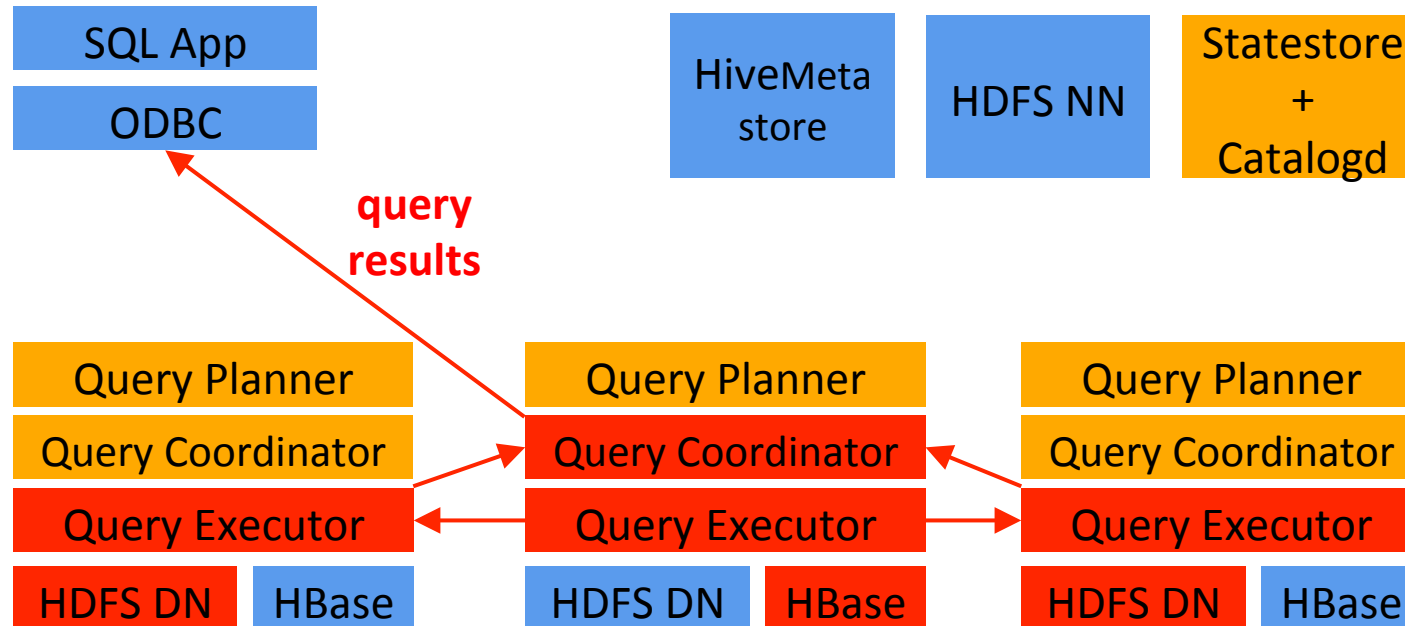
# Impala Architecture: Query Execution

Planner turns request into collections of plan fragments  
Coordinator initiates execution on remote impalad's



# Impala Architecture: Query Execution

Intermediate results are streamed between impalad's Query  
results are streamed back to client



# Query Planning: Overview

---

- 2-phase planning process:
  - single-node plan: left-deep tree of plan operators
  - plan partitioning: partition single-node plan to maximize scan locality, minimize data movement
- Parallelization of operators:
  - All query operators are fully distributed

# Query Planning: Single-Node Plan

---

- Plan operators: Scan, HashJoin, HashAggregation, Union, TopN, Exchange


# Single-Node Plan: Example Query

---

```
SELECT t1.custid,  
       SUM(t2.revenue) AS revenue  
FROM LargeHdfsTable t1  
JOIN LargeHdfsTable t2 ON (t1.id1 = t2.id)  
JOIN SmallHbaseTable t3 ON (t1.id2 = t3.id)  
WHERE t3.category = 'Online'  
GROUP BY t1.custid  
ORDER BY revenue DESC LIMIT 10;
```

# Query Planning: Single-Node Plan

- Single-node plan for example:





# Query Planning: Distributed Plans

---

- Goals:
  - maximize scan locality, minimize data movement
  - full distribution of all query operators (where semantically correct)
- Parallel joins:
  - broadcast join: join is colocated with left input; right-hand side table is broadcast to each node executing join  
-> preferred for small right-hand side input
  - partitioned join: both tables are hash-partitioned on join columns  
-> preferred for large joins
  - cost-based decision based on column stats/estimated cost of data transfers

# Query Planning: Distributed Plans

---

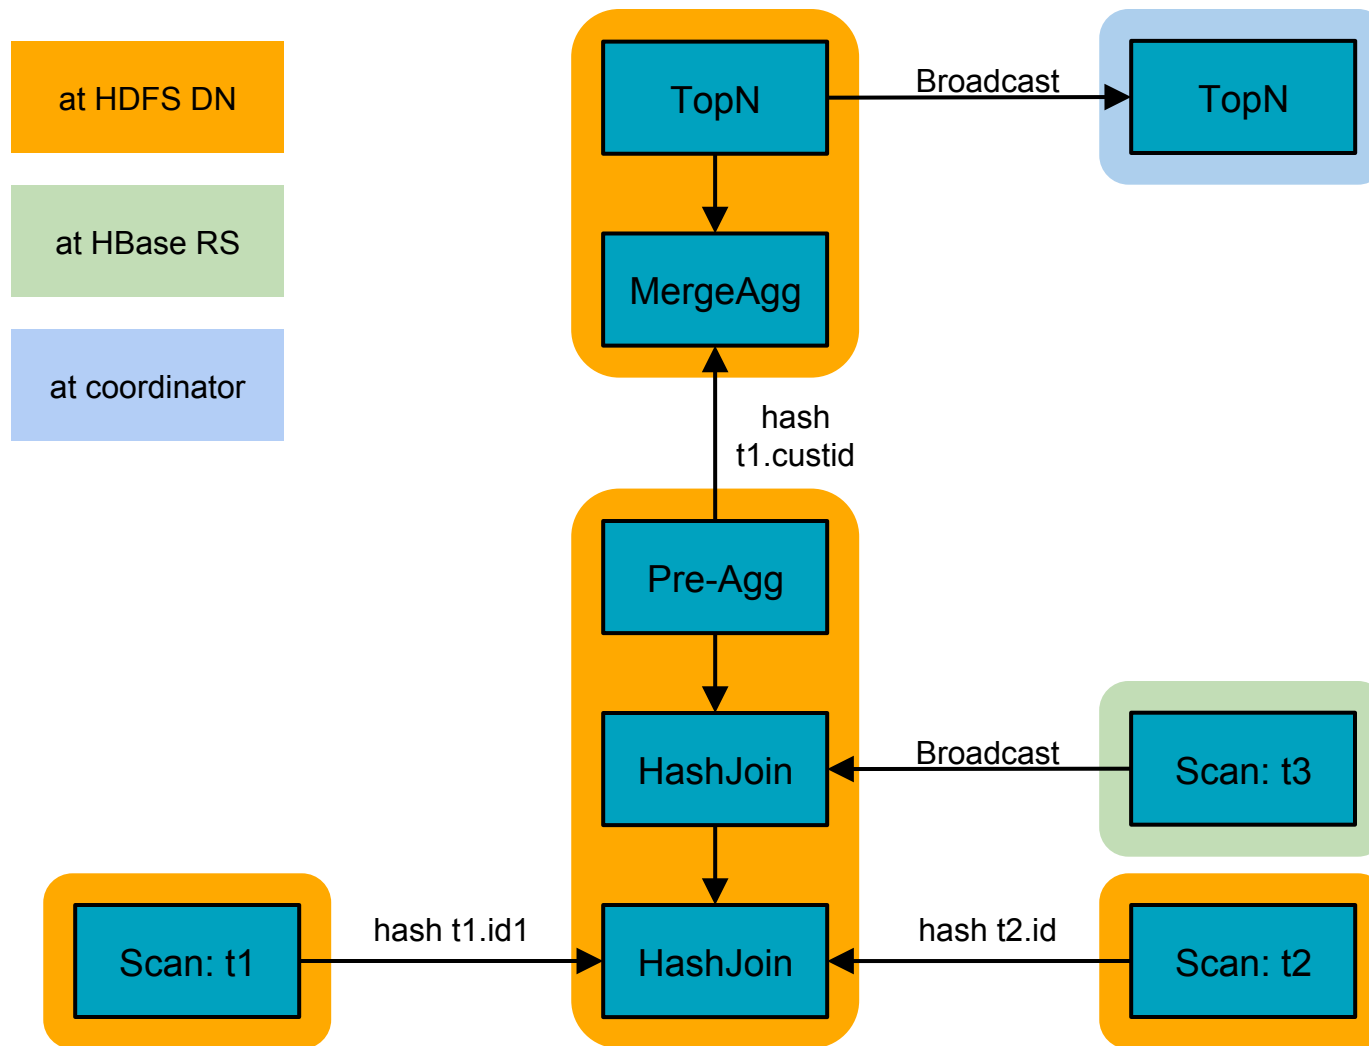
- Parallel aggregation:
  - pre-aggregation where data is first materialized
  - merge aggregation partitioned by grouping columns
- Parallel top-N:
  - initial top-N operation where data is first materialized
  - final top-N in single-node plan fragment

# Query Planning: Distributed Plans

---

- In the example:
  - scans are local: each scan receives its own fragment
  - 1st join: large x large -> partitioned join
  - 2nd scan: large x small -> broadcast join
  - pre-aggregation in fragment that materializes join result
  - merge aggregation after repartitioning on grouping column
  - initial top-N in fragment that does merge aggregation
  - final top-N in coordinator fragment

# Query Planning: Distributed Plans



# Metadata Handling

---

- Impala metadata:
  - Hive's metastore: logical metadata (table definitions, columns, CREATE TABLE parameters)
  - HDFS NameNode: directory contents and block replica locations
  - HDFS DataNode: block replicas' volume ids

# Metadata Handling

---

- Caches metadata: no synchronous metastore API calls during query execution
- impalad instances read metadata from metastore at startup
- Catalog Service relays metadata when you run DDL or update metadata on one of Impalad's
- `REFRESH [ <tbl> ]`: reloads metadata on all impalad's (if you added new files via Hive)
- `INVALIDATE METADATA`: reloads metadata for all tables
- [Roadmap: HCatalog](#)

# Impala Execution Engine

---

- Written in C++ for minimal execution overhead
- Internal in-memory tuple format puts fixed-width data at fixed offsets
- Uses intrinsics/special cpu instructions for text parsing, crc32 computation, etc.
- Runtime code generation for “big loops”

# Impala Execution Engine

---

- More on runtime code generation
  - example of "big loop": insert batch of rows into hash table
  - known at query compile time: # of tuples in a batch, tuple layout, column types, etc.
  - generate at compile time: unrolled loop that inlines all function calls, contains no dead code, minimizes branches
  - code generated using llvm



# Impala's Statestore

---

- Central system state repository
  - name service (membership)
  - Metadata
  - Roadmap: other scheduling-relevant or diagnostic state
- Soft-state
  - all data can be reconstructed from the rest of the system
  - cluster continues to function when statestore fails, but per-node state becomes increasingly stale
- Sends periodic heartbeats
  - pushes new data
  - checks for liveness

# Statestore: Why not ZooKeeper?

---

- ZK is not a good pub-sub system
  - Watch API is awkward and requires a lot of client logic
  - multiple round-trips required to get data for changes to node's children
  - push model is more natural for our use case
- Don't need all the guarantees ZK provides:
  - serializability
  - persistence
  - prefer to avoid complexity where possible
- ZK is bad at the things we care about and good at the things we don't

# Comparing Impala to Dremel

---

- What is Dremel?
  - columnar storage for data with nested structures
  - distributed scalable aggregation on top of that
- Columnar storage in Hadoop: Parquet
  - stores data in appropriate native/binary types
  - can also store nested structures similar to Dremel's ColumnIO
- Distributed aggregation: Impala
- Impala plus Parquet: a superset of the published version of Dremel (which didn't support joins)

# More about Parquet

---

- What is it:
  - container format for all popular serialization formats: Avro, Thrift, Protocol Buffers
  - Successor to Trevn
  - jointly developed between Cloudera and Twitter
  - open source; hosted on github
- Features
  - rowgroup format: file contains multiple horiz. slices
  - supports storing each column in separate file
  - supports fully shredded nested data; repetition and definition levels similar to Dremel's ColumnIO
  - column values stored in native types (bool, int<x>, float, double, byte array)
  - support for index pages for fast lookup
  - extensible value encodings

# Comparing Impala to Hive

---

- Hive: MapReduce as an execution engine
  - High latency, low throughput queries
  - Fault-tolerance model based on MapReduce's on-disk checkpointing; materializes all intermediate results
  - Java runtime allows for easy late-binding of functionality: file formats and UDFs.
  - Extensive layering imposes high runtime overhead
- Impala:
  - direct, process-to-process data exchange
  - no fault tolerance
  - an execution engine designed for low runtime overhead

# Impala Roadmap: 2013

---

- Additional SQL:
  - ORDER BY without LIMIT
  - Analytic window functions
  - support for structured data types
- Improved HBase support:
  - composite keys, complex types in columns, index nested-loop joins, INSERT/UPDATE/DELETE

# Impala Roadmap: 2013

---

- Runtime optimizations:
  - straggler handling
  - improved cache management
  - data collocation for improved join performance
- Resource management:
  - goal: run exploratory and production workloads in same cluster, against same data, w/o impacting production jobs

# Demo

---

- Uses Cloudera's Quickstart VM  
<http://tiny.cloudera.com/quick-start>



# Try it out!

---

- Open source! Available at [cloudera.com](http://cloudera.com), AWS EMR!
- We have packages for:
- RHEL 5,6, SLES11, Ubuntu Lucid, Maverick, Precise, Debain, etc.
- Questions/comments? [community.cloudera.com](http://community.cloudera.com)
- My twitter handle: [mark\\_grover](https://twitter.com/mark_grover)
- Slides at: [github.com/markgrover/impala-thug](https://github.com/markgrover/impala-thug)