

# Adaptive Execution in Baidu

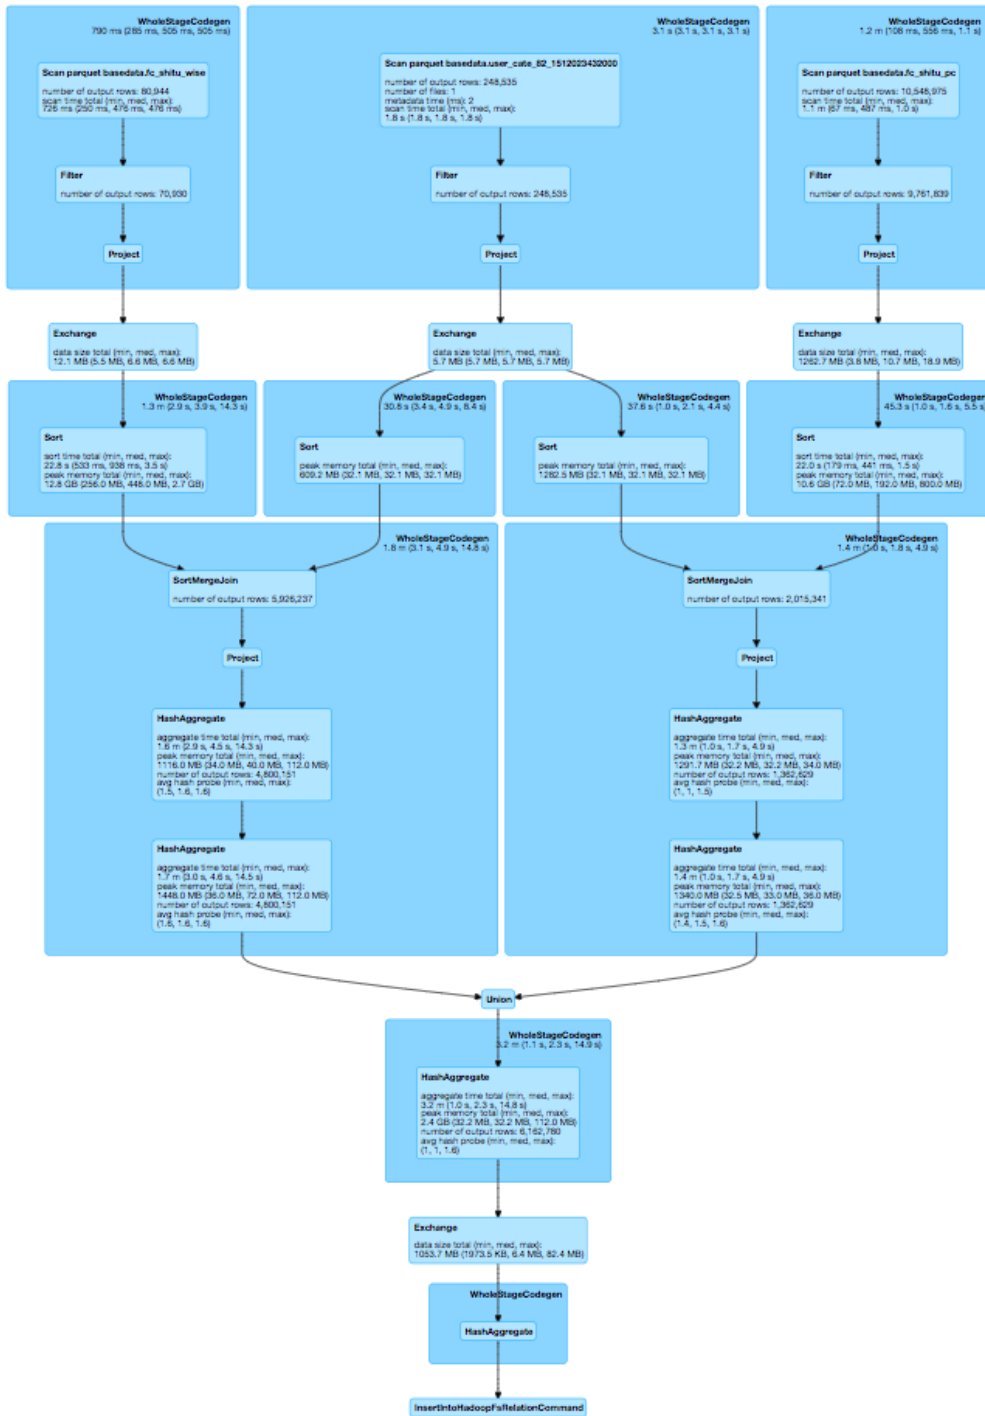
- 1. SortMergeJoin to BroadcastJoin
  - a) with AE closed
  - b) with AE opened
- 2. Long running application or use Spark as a service
- 3. GraphFrame jobs

This document shares some typical user cases about Spark Adaptive Execution in Baidu. Here we divide them mainly into 3 scenario:

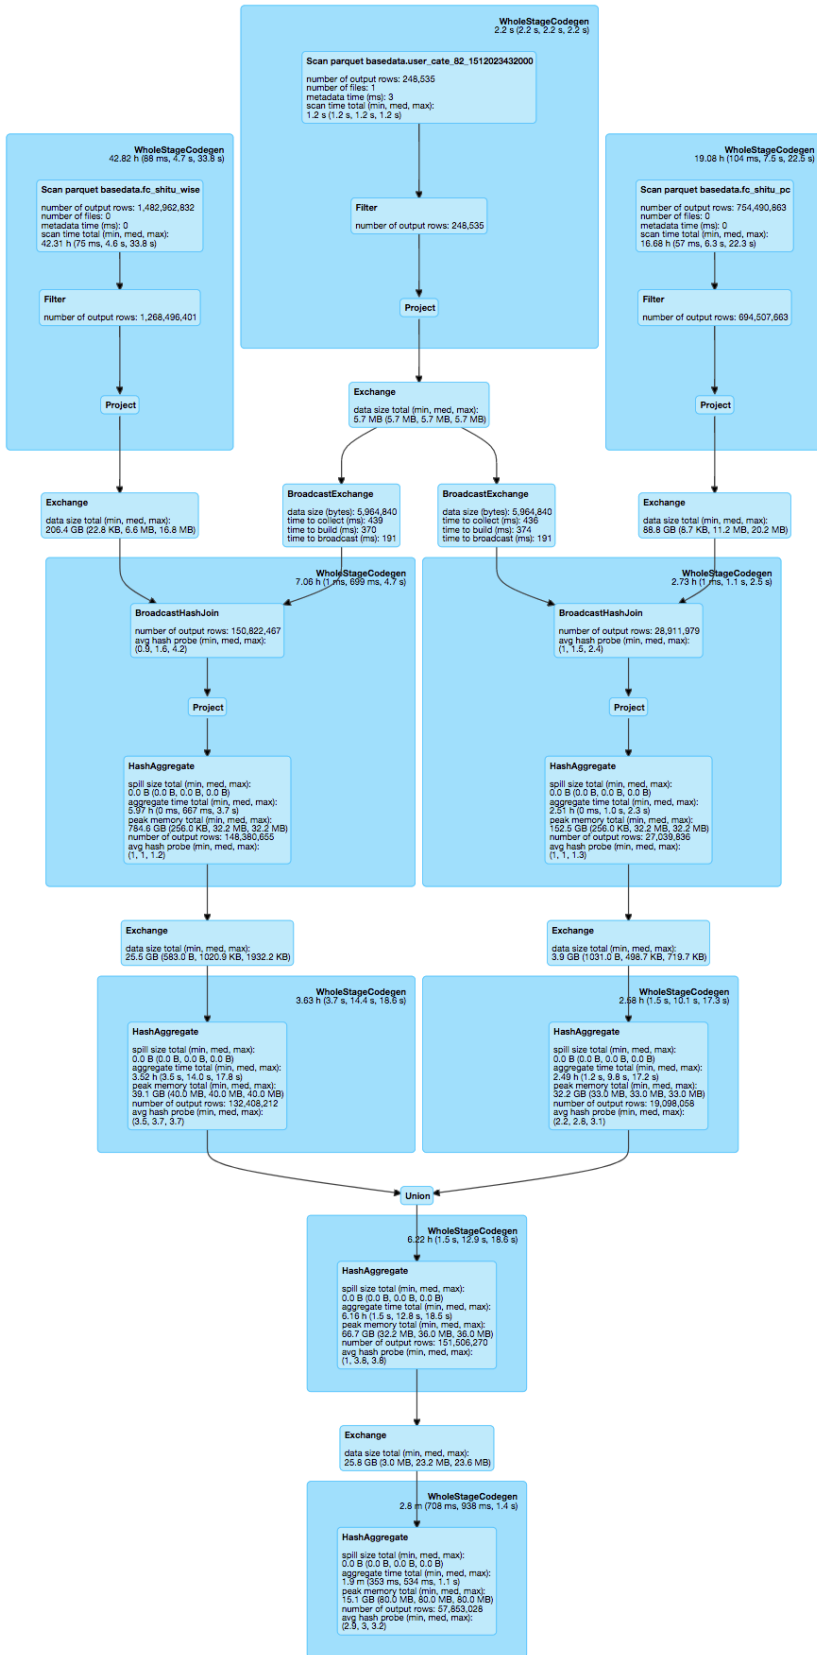
## 1. SortMergeJoin to BroadcastJoin

The SortMergeJoin transform to BroadcastJoin over deeply tree node can bring us **50% to 200%** boosting on query performance, and this strategy always hit the BI scenario like join several tables with filter strategy in subquery, see the example below:

a) with AE closed



b) with AE opened



We can see the both side of SortMergeJoin transform to BroadcastHashJoin, there's no complex config on this scenario, just reused the broadcast threshold. In this case, AE brings us 200%(9min→3min) boosting on performance improvement.

## 2. Long running application or use Spark as a service

In this case, long running application refers to the duration of application near 1 hour. Using Spark as a service refers to use spark-shell and keep submit sql or use the service of Spark like Zeppelin, Livy or our internal sql service Baidu BigSQL. In such scenario, all spark jobs share same partition number, so enable AE and add configs about expected task info including data size, row number, min\max partition number and etc, will bring us 50%-100%boosting on performance improvement.

AE Enable	Application Detail  (AE will generate extra job to do some metrics collection)	Duration	AE config	Resource Config																																										
False	<div>Completed jobs (6)</div> <table><tr><th>Job ID</th><th>Application</th><th>Submitted</th><th>Duration</th><th>Report Submitted (min)</th><th>Task (out of 10000) Submitted (min)</th></tr><tr><td>1</td><td>Spark SQL (hive) [hive-metastore]</td><td>2019-01-10 14:10:10</td><td>51 min</td><td>10 min</td><td>10 min</td></tr><tr><td>2</td><td>Spark SQL (hive) [hive-metastore]</td><td>2019-01-10 14:10:10</td><td>51 min</td><td>10 min</td><td>10 min</td></tr><tr><td>3</td><td>Spark SQL (hive) [hive-metastore]</td><td>2019-01-10 14:10:10</td><td>51 min</td><td>10 min</td><td>10 min</td></tr><tr><td>4</td><td>Spark SQL (hive) [hive-metastore]</td><td>2019-01-10 14:10:10</td><td>51 min</td><td>10 min</td><td>10 min</td></tr><tr><td>5</td><td>Spark SQL (hive) [hive-metastore]</td><td>2019-01-10 14:10:10</td><td>51 min</td><td>10 min</td><td>10 min</td></tr><tr><td>6</td><td>Spark SQL (hive) [hive-metastore]</td><td>2019-01-10 14:10:10</td><td>51 min</td><td>10 min</td><td>10 min</td></tr></table>	Job ID	Application	Submitted	Duration	Report Submitted (min)	Task (out of 10000) Submitted (min)	1	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min	2	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min	3	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min	4	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min	5	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min	6	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min	52 min	null	<div>spark.executor.instances 100</div> <div>spark.executor.memory 10G</div> <div>spark.executor.cores 4</div> <div>spark.task.cpus 1</div>
Job ID	Application	Submitted	Duration	Report Submitted (min)	Task (out of 10000) Submitted (min)																																									
1	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min																																									
2	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min																																									
3	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min																																									
4	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min																																									
5	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min																																									
6	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min																																									
True	<div>Jobs (6)</div> <table><tr><th>Job ID</th><th>Application</th><th>Submitted</th><th>Duration</th><th>Report Submitted (min)</th><th>Task (out of 10000) Submitted (min)</th></tr><tr><td>1</td><td>Spark SQL (hive) [hive-metastore]</td><td>2019-01-10 14:10:10</td><td>51 min</td><td>10 min</td><td>10 min</td></tr><tr><td>2</td><td>Spark SQL (hive) [hive-metastore]</td><td>2019-01-10 14:10:10</td><td>51 min</td><td>10 min</td><td>10 min</td></tr><tr><td>3</td><td>Spark SQL (hive) [hive-metastore]</td><td>2019-01-10 14:10:10</td><td>51 min</td><td>10 min</td><td>10 min</td></tr><tr><td>4</td><td>Spark SQL (hive) [hive-metastore]</td><td>2019-01-10 14:10:10</td><td>51 min</td><td>10 min</td><td>10 min</td></tr><tr><td>5</td><td>Spark SQL (hive) [hive-metastore]</td><td>2019-01-10 14:10:10</td><td>51 min</td><td>10 min</td><td>10 min</td></tr><tr><td>6</td><td>Spark SQL (hive) [hive-metastore]</td><td>2019-01-10 14:10:10</td><td>51 min</td><td>10 min</td><td>10 min</td></tr></table>	Job ID	Application	Submitted	Duration	Report Submitted (min)	Task (out of 10000) Submitted (min)	1	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min	2	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min	3	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min	4	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min	5	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min	6	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min	30 min	<div>spark.sql.adaptive.enabled true</div> <div>spark.sql.adaptive.minNumPostShufflePartitions 400</div> <div>spark.sql.adaptive.maxNumPostShufflePartitions 10000</div> <div>spark.sql.adaptive.shuffle.targetPostShuffleInputSize 512M</div>	<div>spark.executor.instances 100</div> <div>spark.executor.memory 10G</div> <div>spark.executor.cores 4</div> <div>spark.task.cpus 1</div>
Job ID	Application	Submitted	Duration	Report Submitted (min)	Task (out of 10000) Submitted (min)																																									
1	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min																																									
2	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min																																									
3	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min																																									
4	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min																																									
5	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min																																									
6	Spark SQL (hive) [hive-metastore]	2019-01-10 14:10:10	51 min	10 min	10 min																																									

### 3. GraphFrame jobs

The last scenario is the application use GraphFrame, in this case, user has a 2-dimension graph with 1 billion edges, use the connected componentsalgorithm in GraphFrame. With enabling AE, the duration of app reduce from 58min to 32min, almost 100% boosting on performance improvement.